Draft version April 11, 2026
Typeset using LaTeX **twocolumn** style in AASTeX7.0.1

Self-Sustained Penrose Excitation as a Testable Phenomenological Protocol for High-Spin Quasars

Jun Wakabayashi[1]

[1] *Independent Researcher, Japan*

ABSTRACT

We present a high-spin, near-threshold equatorial radiative channel that (i) sustains super-Eddington luminosities and (ii) self-consistently enforces a sub-unity spin ceiling, explaining why spins do not exceed unity in practice.

Compactness-limited transparency (ℓ) gates the observable power, and a simple torque balance links the luminosity to the same control that sets the ceiling.

We provide an executable, screening-based observational protocol; operationally, it predicts three same-epoch hooks at high spin and high accretion rate:
(1) A 20–120 keV high-energy shoulder favored over a featureless continuum by ΔBIC (ΔAICc *optional*) ≥ 6;
(2) Lag hardening in energy-resolved X-ray timing, i.e., hard lags with $\tau(E)$ increasing with energy relative to a fixed soft reference band; and
(3) Equator-aligned X-ray polarization of the hard component with a rising degree (with X-ray→UV/optical reprocessing lags as a complementary Tier-2 check).

The Tier-1 archival test is explicit: after screening for heavy absorption, extreme inclination, and high compactness, recovery of the joint-hook pattern in more than one source supports the existence-style claim, whereas zero or one such source leaves it not established.

Speculative note: We also note that rare geometries may allow partial direct leakage of the quasi-beam, potentially producing fast-variability events (speculative; see Sec. 5.5).

Keywords: Quasars — Radio-quiet quasars, Accretion, accretion disks — Black hole physics — Relativistic processes

1. INTRODUCTION

Summary.— This paper develops a concrete, spin-regulated mechanism—Self-Sustained Penrose Excitation (SSPE)—that can both power super-Eddington quasar luminosities and enforce a near-extremal spin ceiling, while remaining operationally testable by X-ray and polarimetric diagnostics.

Accreting supermassive black holes (SMBHs) display a remarkable range of radiative efficiencies and Eddington ratios.

Standard thin-disk theory struggles to account for quasars radiating persistently at factors of a few above the Eddington limit without invoking either extremely fine-tuned parameters or poorly constrained "slim" accretion flows (M. A. Abramowicz et al. 1988).

At the same time, near-extremal Kerr spins appear to be common among luminous active galactic nuclei (AGN), yet the dimensionless spin parameter a_* is expected to remain bounded below unity by basic causal and thermodynamic arguments.

Any viable model for the most luminous quasars must therefore accomplish two tasks at once: (i) provide an additional, repeatable energy source that can sustain $L \sim 2$–$3\,L_{\rm Edd}$ (with occasional excursions to 5–10 $L_{\rm Edd}$), and (ii) explain why the spin distribution shows an apparent "ceiling" just short of $a_* = 1$.

A second, independent tension concerns the outer thermal structure of luminous accretion disks (E. M. Cackett et al. 2021).

In several objects, the continuum-emitting disk appears more extended and hotter at large radii than

Corresponding author: Jun Wakabayashi

expected from standard thin-disk expectations (C. W. Morgan et al. 2010; Y.-P. Li et al. 2019).

In other words, the disk seems to be "puffed up and overheated" on scales where standard radiative and viscous heating would predict a cooler, geometrically thinner structure.

This raises the possibility that part of the central power budget is redistributed into the equatorial flow.

This combination — a substantial central power budget and a conspicuously hot, inflated outer disk — suggests that a non-negligible fraction of the black-hole spin energy might be being redirected back into the equatorial flow.

Tensions across probes: local fixes vs. global synthesis.— Taken in isolation, each of several long-running tensions can be accommodated by source-by-source adjustments within existing frameworks:

(i) hard X/MeV shoulders can be fit with multi-component Comptonization or corona-geometry variants;
(ii) wavelength-dependent lag spectra can be absorbed into more complex reprocessing geometries;
(iii) microlensing and reverberation sizes can be softened by modest disk thickening and radiative-transfer complications.

However, the recurring appearance of these tensions *in the same high-spin, high-λ population* suggests that treating them as unrelated, probe-specific anomalies is an increasingly high-dimensional patchwork.

This paper frames SSPE as a single *equatorial inner-disk heating / power-transfer channel* that can be confronted with data across these disparate probes under a shared screening logic (compactness, absorption, inclination).

The intent is not to declare uniqueness, but to provide a unifying, operationally testable target that invites coordinated scrutiny by the spectral–timing, microlensing/reverberation, and GRMHD/plasma communities.

This paper introduces and develops *Self-Sustained Penrose Excitation* (SSPE): a quasi-equatorial, Penrose-like energy-transfer channel that couples the ergoregion to the innermost accretion flow.

When the Kerr spin exceeds a near-threshold value $a_* \gtrsim a_{\rm th} \simeq 0.97$, a modest fractional coupling of the rotational-energy reservoir, parameterized by $\epsilon_{\rm coup}(a_*) \sim 10^{-2}$–$10^{-1}$, can continuously re-inject power into the inner disk while providing a counter-torque that self-limits the spin.

Our goal is to establish SSPE as a quantitatively specified, operationally testable framework.

First, we formalize the geometry and kinematics of the ISCO–ergoregion interface and introduce a minimal parameterization of the coupling efficiency and threshold (Section 2).

Second, we present a semi-analytic parameter study exploring how $L/L_{\rm Edd}$, duty cycle, and spin evolution respond to changes in $\epsilon_{\rm coup}$, coupling scale, and effective screening of high-ℓ influx (Section 2.10).

Third, we translate each observational hook into concrete signal-to-noise, model-selection, and polarimetric requirements, and state an explicit screened-domain existence criterion (Section 3).

Fourth, we apply the framework to BZ-favorable, jet-dominated quasars with anomalously hot/extended disks, arguing that equatorial feedback can inflate and overheat the outer disk even in such systems (Section 5.3).

We do not attempt to replace super-Eddington disk models (e.g., slim disks); rather, SSPE adds a spin-regulated, testable protocol layer whose discriminant is the joint appearance of Hooks 1–3 and the presence of a sub-unity spin ceiling.

Scope and modeling stance.—We present a phenomenological framework: the microphysics of beam formation, transport, and equatorial deposition is intentionally kept agnostic, while the *observable consequences* of such coupling are made explicit and operationally testable.

This stance allows clean confrontation with data now and provides a clear interface to future GRMHD/plasma studies that could instantiate the coupling agent.

Throughout we deliberately restrict attention to a single Kerr SMBH and its immediate accretion environment; issues of cosmological SMBH growth, seeding, and host-galaxy evolution are deferred to separate work.

Boundary conditions are not fixed near $a_* \to 1$.—Much of the literature treats the inner disk, magnetosphere, and ergoregion with effectively fixed cross-component boundary conditions.

We instead emphasize that near-extremal spin the boundaries themselves evolve: frame dragging enlarges and reshapes the ergoregion, the ISCO approaches it, and the equatorial return path becomes thin and resistive.

In this coupled regime an equatorial current sheet is generically required by the global field topology; shear and flux loading drive plasmoid-dominated reconnection; and the resulting split redistributes energy and angular momentum so that a negative-energy branch is absorbed by the hole while a positive-energy branch vents as a narrow equatorial quasi-beam, part of which returns to heat the inner disk.

2. THEORETICAL FRAMEWORK

Summary.— Near-threshold Kerr coupling at the ISCO–ergoregion interface supplies disk heating while extracting spin.

Its cleanest observable imprint is expected within a compactness-limited transparent window $\ell < 30$.

All elements invoked below act outside the event horizon; the coupling operates in the ergoregion and deposits energy in the equatorial flow, avoiding assumptions about interior or singularity-scale physics.

Abbreviations and references. We use *FFE* for the force-free, magnetically dominated limit ($\rho_e\mathbf{E} + \mathbf{J}\times\mathbf{B} \approx 0$, $\mathbf{J}\cdot\mathbf{E} = 0$) and *GRMHD* for ideal magnetohydrodynamics evolved on a Kerr background (S. S. Komissarov 2004; S. E. Gralla & T. Jacobson 2014; W. E. East & H. Yang 2018; Z. Pan 2018).

We use *PIC* for first-principles kinetic simulations that resolve reconnection and plasmoid formation in the ergoregion current sheet (K. Parfrey et al. 2019; A. Bransgrove et al. 2021).

Throughout we group references as (FFE/GRMHD) versus (PIC).

2.1. *Kerr Energy Reservoir*

For a Kerr black hole of mass M and spin a_*, the extractable rotational energy is

$$E_{\rm rot}(a_*) = \left[1 - \sqrt{\tfrac{1}{2}\left(1 + \sqrt{1 - a_*^2}\right)}\right] Mc^2, \qquad (1)$$

reaching $\sim 0.29\, Mc^2$ as $a_* \to 1$.

2.2. *Spin-Triggered Coupling*

We posit a threshold spin $a_{\rm th}$ above which an ergoregion-coupled agent activates.

The effective coupling follows

$$\epsilon_{\rm coup}(a_*) = \begin{cases} 0, & a_* \le a_{\rm th}, \\ \epsilon_{\rm max}\left(\dfrac{a_* - a_{\rm th}}{1 - a_{\rm th}}\right)^n, & a_* > a_{\rm th}, \end{cases} \qquad (2)$$

where $\epsilon_{\rm max}$ is the saturation cap and n controls activation sharpness.

Activation threshold. The equatorial coupling turns on steeply once the spin exceeds a threshold $a_{\rm th}$.

In our fiducial calibration, the rise is rapid near $a_* \simeq 0.97$, setting the lever for both luminosity boost and spin regulation (cf. Eqs. (3)–(11)).

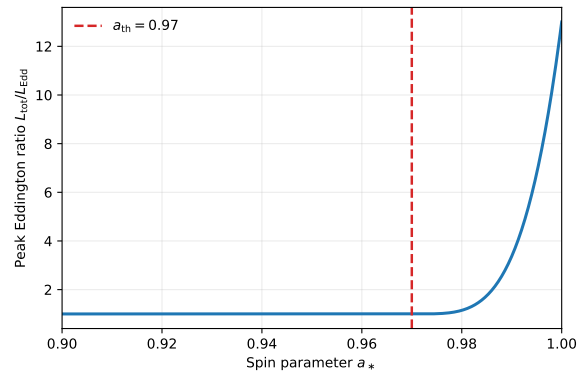

Figure 1. Activation turns on steeply above the threshold. Peak $L_{\rm tot}/L_{\rm Edd}$ rises rapidly once $a_* > a_{\rm th}$, setting the lever for luminosity and spin regulation (cf. Eqs. (2)–(11)).

We adopt a fiducial $a_{\rm th} \simeq 0.97$ as a pragmatic "near-extremal" boundary where the ISCO approaches the ergosurface boundary in the equatorial plane and small changes in a produce disproportionately large changes in horizon-frame energetics and coupling geometry; the precise value is not fundamental and can be refined empirically.

Unless noted, we adopt the following fiducials for figures and estimates: $M = 10^9 M_\odot$, $\eta_{\rm acc} = 0.1$, $\epsilon_{\rm max} = 0.1$, $n = 2$, $R_0 = 10^2$, and $\alpha = 50$; the illustrative threshold is $a_{\rm th} \simeq 0.97$.

Micro-bridge. While we keep microphysics agnostic, the fiducial range $\epsilon_{\rm coup} \sim 10^{-2}$–$10^{-1}$ is broadly consistent with energy-release fractions seen in near-horizon, plasmoid-dominated reconnection in recent GRMHD studies.

For the present framework we only assume that $\epsilon_{\rm coup}$ rises sharply above $a_{\rm th}$ and saturates below $\epsilon_{\rm max}$; detailed calibration is left for future simulations.

Phenomenological parameters and physical ranges. The cap $\epsilon_{\rm max}$ limits equatorial deposition efficiency by energy-budget and pair-compactness constraints; n controls activation sharpness above $a_{\rm th}$ as an effective criticality index of the coupling geometry.

(R_0, α) regulate the dissipation footprint to maintain transparency, trading compactness against reprocessing.

We restrict these to physically plausible ranges and view them as interfaces for future GRMHD/plasma calibration, not curve-fitting knobs.

Sensitivity and saturation. The activation $\epsilon_{\rm coup}(a_*)$ is intentionally steep: $(a_{\rm th}, n)$ set the narrowness of the "spin ceiling".

In practice $\epsilon_{\rm coup}$ and the leakage fraction depend nonlinearly on flux loading and $\dot{M}$, with possible saturation arbitrarily close to the extremal limit.

We therefore treat $(a_{\rm th}, n, \epsilon_{\rm max}, \alpha)$ as calibration parameters to be fixed by future GRMHD-in-ergoregion testbeds.

Power partition. We allow the extracted rotational power to partition into a polar Blandford–Znajek (BZ) channel, an equatorial feedback channel, and other loss channels:

$$P_{\rm ext} = P_{\rm BZ} + P_{\rm eq} + P_{\rm oth}, \qquad f_{\rm BZ} + f_{\rm eq} + f_{\rm oth} = 1.$$

Near $a_* \to 1$, $f_{\rm eq}$ can become substantial (radio-quiet, radiation-dominated states), while $f_{\rm BZ}$ dominates in radio-loud systems; hybrid states are possible in transient MAD-like regimes.

Qualitative selection rule (geometry/flux). The polar BZ-like channel is expected to dominate when the horizon-threading poloidal flux is large and coherent (i.e., a near-MAD-like state, indicated by strong jet proxies or radio loudness).

By contrast, an equatorial channel becomes comparatively more relevant when the inner flow is thick and time-dependent (high-β / radiation-pressure dominated inner flow; moderate coherent vertical flux) and the accumulated large-scale vertical flux is not extreme, so that dissipation and reconnection concentrate in the equatorial return-current sheet.

In this regime the requirement for SSPE is not a finely tuned nozzle, but simply that a nonzero fraction of reconnection ejecta vents through the equatorial channel ($f_{\rm eq} > 0$), with the magnitude of $f_{\rm eq}$ to be constrained by joint observational coherence rather than by choosing among GRMHD/FFE/PIC closures.

2.3. *Mechanism sketch: a leaky equatorial return path*

More generally, the present framework does not require the equatorial channel to arise exclusively from a strict Penrose split.

In a strongly frame-dragged inner flow, flyby-like energy–angular-momentum exchange may also populate equatorial outflow trajectories, without requiring commitment to a single microphysical realization.

We therefore treat the equatorial channel in a mechanism-agnostic way and test for its observational footprint via the Tier-1 hooks rather than by adjudicating a unique launch mechanism.

Terminology. We refer to the power-budget component $P_{\rm eq}$ (fraction $f_{\rm eq}$) as the *equatorial channel*.

The transport agent of this channel is modeled here as an equatorially concentrated outflow component (an "equatorial quasi-beam" in schematic shorthand), i.e., an intermittently launched positive-energy branch originating near the ISCO–ergoregion interface.

We use *ergoregion* for the physical condition of negative-energy states (Penrose splitting), and *ergosurface* for the geometric boundary (Fig. 2). We avoid 'ergosphere' except as informal shorthand.

We use "Penrose process" in the strict Kerr sense of negative-energy states, and "Penrose-like coupling" as a phenomenological label for ergoregion-mediated energy transfer, capturing the essential energetics of the Penrose process while remaining agnostic.

Why an equatorial channel is plausible (without microphysical adjudication). In rapidly spinning systems an equatorial return-current sheet is expected to form where field polarity reverses across the inner flow; time-dependent reconnection in such sheets generically produces plasmoids and intermittent ejecta in GRMHD/FFE and PIC studies (references in Fig. 2).

Our framework does not derive a nozzle or collimation from first principles.

It only requires that a non-negligible fraction of the reconnection ejecta can vent through an equatorially concentrated branch, whose relevance is to be decided by joint observational coherence (Hooks 1–3) rather than by adjudicating GRMHD/FFE/PIC closure here.

Projectile sufficiency.—Reconnection ejecta (plasmoids) in the ergoregion carry specific energy E and angular momentum L. A Penrose split (i.e., the negative-/positive-energy branch separation) requires a branch with $E - \Omega_{\rm H} L < 0$ (horizon condition). [2] [3]

Tension-driven redistribution during ejection changes L by $\Delta L \sim \mathcal{O}(r_{\rm g} v_\phi)$, and reconnection outflows can reach $v \sim 0.1$–$0.5c$ in high-S sheets; hence tens-of-percent shifts in L can suffice to place a fraction on the negative-energy branch, with the complement forming an equatorially concentrated positive-energy escape branch.

Near-extremal spin brings the ISCO into close contact with the ergoregion and strengthens the coupling between the inner flow and the Kerr magnetosphere.

Rather than deriving collimation from first principles, we adopt a minimal interface motivated by time-dependent studies: equatorially symmetric configurations can develop reconnecting equatorial current sheets that become plasmoid-unstable and launch intermittent outflows (e.g., FFE/GRMHD and PIC studies cited in Fig. 2).

$$E - \Omega_{\rm H} L < 0. \tag{3}$$

See R. Penrose 1969; R. Penrose 2002.

A modest back-flow coupling ($\epsilon_{\rm coup} \sim 10^{-2}$–$10^{-1}$) is then sufficient, at the phenomenological level adopted here, to heat the inner disk and help regenerate magnetic flux, thereby closing a self-sustained loop.

For the present phenomenology, however, we do not require a fully specified microphysical launch prescription.

It is sufficient that a non-negligible subset of reconnection-driven ejecta can access equatorially concentrated positive-energy escape paths, while the complementary branch accounts for the required energy–angular-momentum extraction.

Whether this occurs through a strict Penrose split or a closely related ergoregion-mediated exchange is left to future GRMHD/FFE/PIC calibration.

2.4. *Scope/definitions*

We use "Penrose-like extraction" strictly as a phenomenological label for an energy source that may operate in Kerr systems; it is **not** asserted as a unique or sufficient condition for SSPE.

The present work therefore does not address the activation boundary or microphysics of Penrose extraction.

Instead, the testable claim is observational: in **vetted high-spin/high-λ** systems, the SSPE contribution predicts a joint-hook coherence within the same epoch under the pre-registered protocol (Table 1; Appendix A).

Signatures detected with comparable frequency in low-spin control samples would indicate a spin-insensitive origin (general disk/corona physics) and thus fall outside the specific SSPE interpretation targeted here, without altering the protocol's Tier-1 decision bookkeeping.

[2] Notation follows Fig. 2.

[3] where $\Omega_{\rm H}$ is the horizon angular velocity of a Kerr black hole.

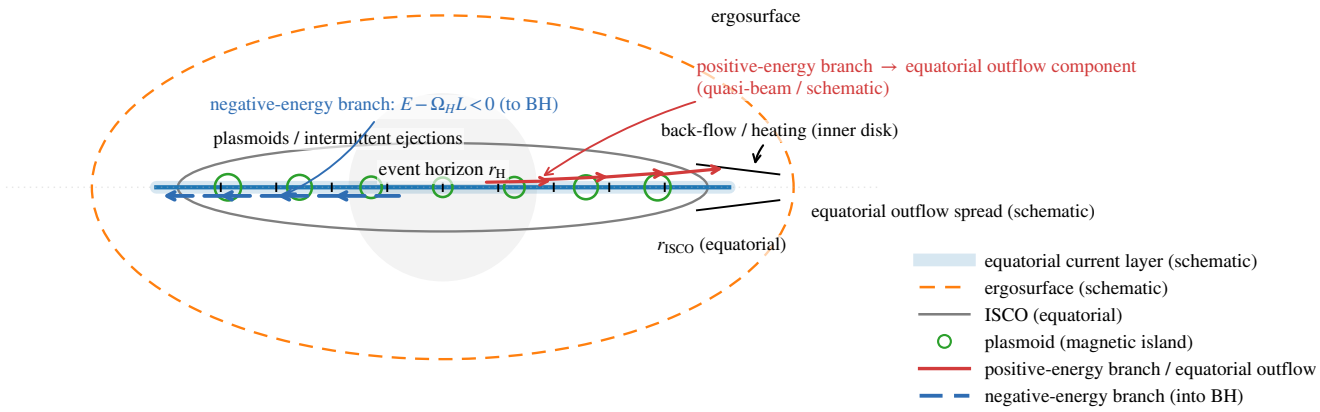

Figure 2. Equatorial ergoregion trigger: negative vs. positive branches (schematic).
Plasma from the inner disk/plunge region enters the ergoregion (inside the ergosurface), where energy–angular-momentum exchange can populate both negative- and positive-energy trajectories/branches.
A **negative-energy branch** with $E - \Omega_H L < 0$ falls through the event horizon (r_H), while a **positive-energy branch** can contribute to an equatorial quasi-beam or, more generally, to an equatorially concentrated outflow that reheats the inner rim near $r_{\rm ISCO}$.
Landmarks—ergosurface, $r_{\rm ISCO}$, and the equatorial return-current sheet (blue ribbon)—are indicated; red/blue branches mark a representative phenomenological realization discussed in the text.
(FFE/GRMHD: S. S. Komissarov 2004; W. E. East & H. Yang 2018; Z. Pan 2018; PIC: K. Parfrey et al. 2019; A. Bransgrove et al. 2021).

Projectile plausibility—Recent XRISM observations of NGC 3783 reported the rapid emergence of an ultrafast outflow synchronized with the decay of a soft X-ray/UV flare, with CME-like kinematics suggesting magnetic driving and possibly reconnection near the disk (L. Gu et al. 2025).

While not an ergoregion-specific probe, this supports the empirical plausibility of intermittent reconnection-driven "projectiles" in inner AGN environments. [4]

Why an equatorial branch?—We do not claim that the equatorial channel must dominate, nor do we fix the microphysical launch/transport details in this paper.

Instead, we treat an intermittently equatorial channel as an operationally testable working condition: if present at a non-negligible level, it can be reprocessed by the inner disk/corona into the joint hooks described in Secs. 3.1–3.

Whether this branch is realized in real quasars is decided by the co-occurrence of Hooks 1–3 under the screens, not by microphysical closure here.

The *positive-energy branch* may therefore propagate as an equatorially concentrated outflow component rather than escaping exclusively along the poles.

Because the transport remains largely collisionless/Poynting-dominated until it reprocesses in the inner disk, and because the dissipation footprint is extended ($R_{\rm eff} \sim 10^2$–$10^3\, r_{\rm g}$), the associated pair compactness can remain low, consistent with the transparency requirement of Eq. (6).

[4] On the XRISM report on NGC 3783, the empirical plausibility of intermittent "projectiles" was stated as an explicit working assumption in the OSF pre-registered protocol for this work (DOI:10.17605/osf.io/q7dth).

2.5. *Self-Sustained Penrose Excitation*

The mean extraction power is $\langle P_{\rm ext}\rangle = E_{\rm rot}/\tau$, and with duty cycle d the instantaneous power during active phases is $P_{\rm ext} \sim \langle P_{\rm ext}\rangle/d$. Here τ denotes the effective extraction e-folding timescale of the rotational-energy reservoir, i.e., $\langle P_{\rm ext}\rangle = E_{\rm rot}/\tau$ using Eq. (1).

2.6. *Disk Dissipation and Scale*

A fraction $\epsilon_{\rm coup}(a_*)$ of $P_{\rm ext}$ is deposited into the disk:

$$L_{\rm self} = \epsilon_{\rm coup}\, P_{\rm ext}, \tag{4}$$

and the dissipation spreads over an effective radius

$$R_{\rm eff} \approx R_0\Big[1 + \alpha\Big(\frac{\epsilon_{\rm coup}}{\epsilon_{\rm max}}\Big)\Big]\, r_g, \tag{5}$$

with $r_g = GM/c^2$, fiducial $R_0 \sim 10^2$ and $\alpha \sim 50$.

Self-transparency (physical note).—The scaling $R_{\rm eff} \approx R_0[1+\alpha(\epsilon_{\rm coup}/\epsilon_{\rm max})]\, r_{\rm g}$ phenomenologically captures geometric spreading, multi-zone deposition, and increased scattering mean free paths as the equatorial quasi-beam heats and rarefies the inner disk corona. Our results require $R_{\rm eff} \sim 10^2$–$10^3\, r_{\rm g}$ during bright episodes to keep the pair compactness low (cf. Sec. 2.7).

2.7. *Transparency (Compactness Constraint)*

Transparency requires pair compactness

$$\ell = \frac{L_{\rm self}\sigma_T}{4\pi R_{\rm eff}\, m_e c^3} \lesssim 30, \tag{6}$$

which couples Eqs. (4) and (5) and motivates $R_{\rm eff} \sim 10^2$–$10^3\, r_g$ during bright episodes.

We adopt $\ell \lesssim 30$ as a conservative transparency threshold following classic compactness arguments; the precise value depends on geometry and spectrum and can be re-tuned in data applications.

See, e.g., (R. Svensson 1984; A. P. Lightman & A. A. Zdziarski 1987).

2.8. *Spin and Mass Evolution*

The horizon angular frequency is

$$\Omega_{\rm H} = \frac{a_* c^3}{2GM\left(1+\sqrt{1-a_*^2}\right)}, \tag{7}$$

where $r_H = r_g\left(1+\sqrt{1-a_*^2}\right)$ and $r_g = GM/c^2$. Evolution obeys

$$\frac{dM}{dt} = \frac{dM_{\rm acc}}{dt} - \frac{P_{\rm ext}}{c^2}, \tag{8}$$

$$\frac{dJ}{dt} = \frac{dJ_{\rm acc}}{dt} - \frac{P_{\rm ext}}{\Omega_{\rm H}}, \tag{9}$$

$$\frac{da_*}{dt} = \frac{c}{GM^2}\frac{dJ}{dt} - 2a_*\frac{1}{M}\frac{dM}{dt}. \tag{10}$$

Torque-balance interpretation of the spin ceiling.—Figure 3 visualizes the torque balance behind the spin ceiling in our minimal evolution model (Eqs. (8)–(10)).

These relations are intended as schematic bookkeeping equations for the minimal spin-budget picture adopted here, rather than as a closed GRMHD torque solution.

For a given coupling strength $\epsilon_{\rm coup}$, we evaluate da_*/dt as the sum of accretion-driven spin-up and SSPE extraction-driven spin-down.

In this setup, a larger $\epsilon_{\rm coup}$ corresponds to a more effective transfer of extracted spin energy (i.e., a stronger counter-torque acting on the hole), which shifts the zero-crossing of da_*/dt to lower a_*.

The resulting equilibrium spin $a_{\rm eq}$ (defined by $da_*/dt = 0$) therefore decreases systematically with increasing $\epsilon_{\rm coup}$.

We define the equilibration timescale $\tau_{\rm eq}$ as the cumulative active time required for a_* to reach $a_{\rm eq}$ within a fixed tolerance; this is the relevant clock when SSPE operates intermittently.

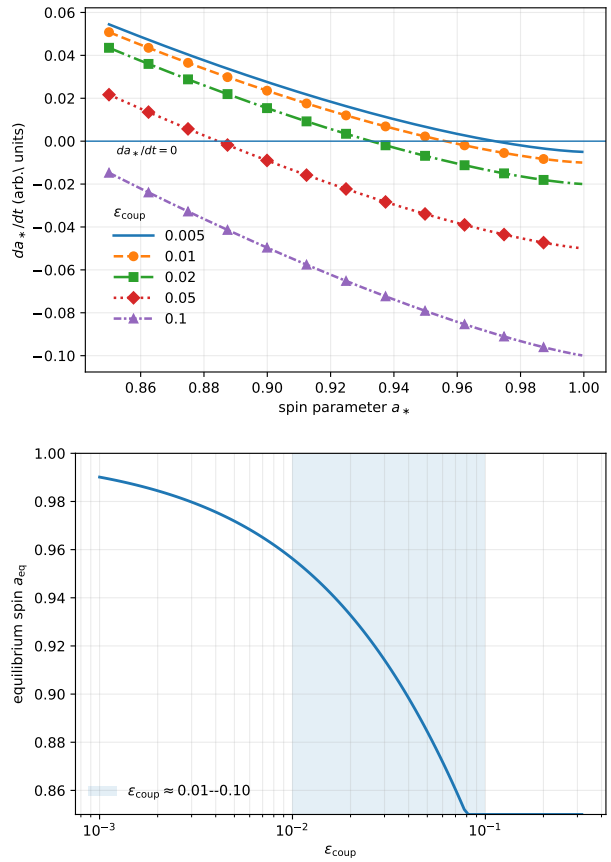

Figure 3. Spin evolution and equilibrium vs. coupling.
Top: Spin drift rate da_*/dt from Eqs. (8)–(10). Negative values imply net spin-down. Normalization is arbitrary: only sign and zero-crossing are used.
Bottom: Equilibrium spin $a_{\rm eq}$ obtained from $da_*/dt = 0$ as a function of $\epsilon_{\rm coup}$; The shaded band indicates the fiducial range, $\epsilon_{\rm coup} \sim 10^{-2}$–$10^{-1}$, used elsewhere. See also Eq. (7).

Spin-down Timescale—Figure 4 illustrates the resulting spin evolution under equatorial extraction for a fiducial coupling, and shows that the effective timescale is set by the integrated high-λ duty cycle rather than a single continuous episode.

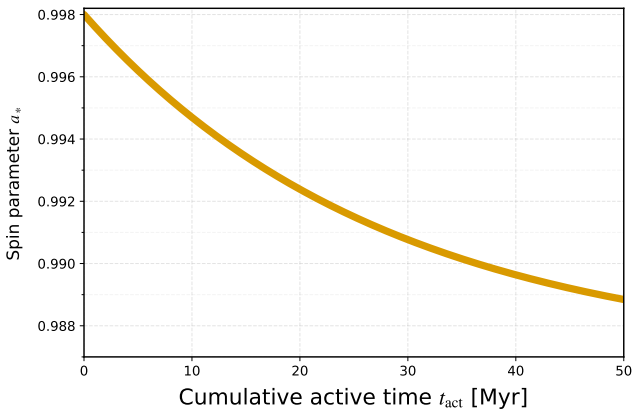

Figure 4. Spin-down on Myr scales under equatorial extraction. Time evolution of the spin parameter $a_*(t)$ obtained by integrating Eqs. (8)–(10) with horizon frequency Eq. (7) for a fiducial coupling $\epsilon_{\rm coup}$ (solid curve, $\epsilon_{\rm coup} = 0.05$). The monotonic drift toward $a_{\rm eq} < 1$ illustrates the self-limiting ceiling. The timescale scales approximately with the cumulative active time; for intermittent high-λ episodes with duty cycle d, the approach to $a_{\rm eq}$ is slowed by $\sim 1/d$ relative to continuous extraction. See Sec. 2.10.

Timescale consistency (active time vs. quasar lifetime).—Figure 4 is plotted against the *cumulative active time* spent in the high-λ SSPE-active state, rather than wall-clock time.

If activity is intermittent with duty cycle d, the corresponding elapsed time is $t_{\rm real} \simeq t_{\rm act}/d$, but the spin evolution depends primarily on $t_{\rm act}$ because extraction is effective only during active episodes.

For the fiducial coupling used in Eqs. (8)–(10), the characteristic drift toward $a_{\rm eq} < 1$ occurs over $t_{\rm act} \sim$ Myr–tens of Myr (Fig. 4), implying $t_{\rm real} \sim (t_{\rm act}/d)$.

For plausible quasar lifetimes of $\sim$ 10–100 Myr and moderate intermittency ($d \sim 0.1$–1), the required cumulative active time can be accumulated within a single luminous phase or across repeated high-λ episodes.

For example, $t_{\rm act} = 20\,$Myr with $d = 0.2$ implies $t_{\rm real} \simeq t_{\rm act}/d \approx 100\,$Myr.

Conversely, in short-lived or highly intermittent systems where $t_{\rm act} \ll \tau_{\rm eq}$, the framework predicts a surviving high-spin tail even at high λ; this is an observationally testable outcome rather than a loophole.

2.9. *Net Luminosity*

The total luminosity is

$$L_{\rm tot} = L_{\rm acc} + L_{\rm self}, \qquad L_{\rm acc} \approx \eta_{\rm acc}\left(\frac{dM_{\rm acc}}{dt}\right)c^2, \quad (11)$$

which links the dynamical solution (Eqs. (8)–(10)) to observables.

Figure 5 summarizes the equilibration time and the mapping to $L_{\rm tot}/L_{\rm Edd}$ at $a_{\rm eq}$.

Energetic sanity check (one-line).—With $M_{\rm BH} = 10^9\,M_\odot$ [$Mc^2 \simeq 1.8 \times 10^{63}$ erg] and $E_{\rm rot} \sim 0.1\,Mc^2$, a reservoir e-fold $\tau = 10^7$ yr, $\epsilon_{\rm coup} = 0.05$, and duty $d = 0.2$ yield $L_{\rm self} \sim (E_{\rm rot}/\tau)\,(\epsilon_{\rm coup}/d) \approx 1.4 \times 10^{47}$ erg s^{-1} $\sim 1.1\,L_{\rm Edd}$, and 2–3 $L_{\rm Edd}$ when combined with concurrent accretion, while $da_*/dt < 0$ in this minimal picture prevents overspin.

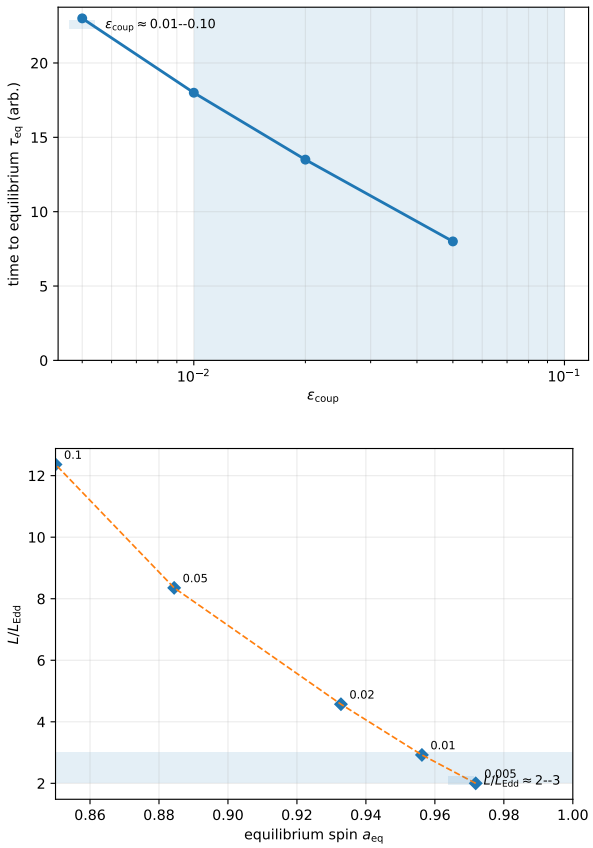

Figure 5. Equilibrium timescale and luminosity mapping.
Top: Time to equilibrium $\tau_{\rm eq}$ versus coupling $\epsilon_{\rm coup}$, computed from the spin-evolution system Eqs. (8)–(10) with horizon frequency Eq. (7); the shaded band $\epsilon_{\rm coup} \sim 10^{-2}$–$10^{-1}$ is the fiducial range used elsewhere. We define the equilibration timescale $\tau_{\rm eq}$ as the cumulative active time required to reach $|a_* - a_{\rm eq}| < 10^{-4}$; the absolute normalization is set by the integration step, so we use only the scaling with $\epsilon_{\rm coup}$.
Bottom: Total luminosity at equilibrium, $L_{\rm tot}/L_{\rm Edd}$ from Eq. (11) evaluated at $a_* = a_{\rm eq}$, showing a typical 2–3 band (blue) with labeled $\epsilon_{\rm coup}$ values along the trend.

2.10. *Semi-analytic parameter study*

Analytic Exploration—The semi-analytic exploration above yields three robust trends that do not depend on the microphysical details of how the equatorial channel is realized.

First, once the coupling activates above the threshold spin $a_* \gtrsim a_{\rm th}$, even a modest efficiency $\epsilon_{\rm coup} \sim 10^{-2}$–$10^{-1}$ can supply an inner-disk feedback luminosity $L_{\rm self} = \epsilon_{\rm coup} P_{\rm ext}$ that is comparable to $L_{\rm Edd}$ for plausible reservoir and duty-cycle choices, so that the combined output $L_{\rm tot} = L_{\rm acc} + L_{\rm self}$ generically reaches the observed $L/L_{\rm Edd} \sim 2$–3 band.

Second, the same extraction term necessarily carries angular momentum and therefore introduces a counter-torque; integrating the minimal evolution system generically drives a_* toward an equilibrium $a_{\rm eq} < 1$ on Myr-scale equilibration times, with stronger coupling yielding a lower $a_{\rm eq}$.

Third, compactness links the observability of the hard component to geometry: the requirement $\ell \lesssim 30$ couples $(\epsilon_{\rm coup}, d)$ to the dissipation footprint $R_{\rm eff}$, carving a transparency-limited window in which bright states remain pair-safe.

In practice, luminous high-λ phases may be intermittent.

Accordingly, the relevant condition for spin regulation is that the *integrated* time spent in SSPE-active states (above $a_{\rm th}$ and within the compactness-safe window) is comparable to or exceeds the effective equilibration time.

If vetted high-spin/high-λ sources exhibit no spin ceiling despite sufficient cumulative active time, the framework would be disfavored.

Trends Map—These trends map directly onto the observational program in Sec. 3.

In the compactness-safe window (i.e., after screening out high-ℓ states), SSPE predicts that the hard component associated with equatorial feedback should be simultaneously (i) spectrally visible as a high-energy shoulder, (ii) temporally imprinted via lag-energy hardening consistent with inner-disk reheating and reprocessing, and (iii) geometrically encoded through disk-parallel hard X-ray polarization.

Because activation is steep near $a_{\rm th}$, the hooks are expected, when present, to concentrate in the high-spin, high-λ subset rather than appearing as unrelated, source-by-source anomalies.

Conversely, in short-lived or highly intermittent systems where $t_{\rm act} \ll \tau_{\rm eq}$, the framework predicts a surviving high-spin tail even at high λ; this is an observationally testable outcome rather than a loophole.

Finally, the same bookkeeping clarifies why SSPE should be treated as an equatorial complement to polar BZ extraction: across states, a changing partition between $P_{\rm BZ}$ and $P_{\rm eq}$ can produce an anti-correlation between jet dominance and equatorial reheating diagnostics, with rare geometries potentially allowing partial direct leakage discussed separately in Sec. 5.5.

SECTION SUMMARY AND PROTOCOL ROADMAP

1. **Semi-analytic picture (this section):** Once the coupling activates above a threshold spin, an equatorial radiative channel can sustain $L/L_{\rm Edd} > 1$ while naturally driving the spin toward an equilibrium ceiling.
2. **What sets the state:** The channel is powered by repeated Penrose-like coupling at the ISCO–ergoregion interface, depositing power back into the inner disk while extracting angular momentum.
3. **Why observability is gated:** Compactness links detectability to geometry; a transparency-limited window explains why the hard component is seen only in specific bright states.
4. **Next (Sec. 3):** We translate the framework into an operational screening protocol with minimal Screens (A–D) and three joint hooks (spectral, timing, polarimetric), together with explicit per-hook adjudication criteria (Table 1).
5. **Tier-1 archival test (Sec. 5):** In vetted high-spin/high-λ samples, recovery of the SSPE-like joint-hook pattern in more than one vetted source supports the existence-style claim. If zero or one such source remains, the claim is not established; if screened or decidable coverage is too limited, the archival result is treated as inconclusive.

3. OBSERVATIONAL HOOKS AND SCREENING PROTOCOL

Summary.— *Screens:* select high-λ systems with favorable geometry/obscuration and compactness.
Hooks: search for a hard X/MeV "shoulder," lag hardening with energy, and equator-aligned polarization signatures.
Tier-1 archival test: if the vetted high-spin/high-λ archival set does not recover the joint-hook pattern in more than one source, the existence-style claim is not established in that regime.

In this paper we treat SSPE as a concrete, operationally testable working framework rather than a completed theory.

This work does not adjudicate between GRMHD/FFE/PIC implementations; it treats an equatorial channel as an operationally testable working condition and targets observational joint coherence.

The microphysics of the coupling between the ergoregion, quasi-equatorial channels, and the inner disk is intentionally left agnostic and absorbed into a small set of phenomenological parameters.

What is fixed — and offered to observers and simulators — is an executable protocol: (i) minimal screens that define vetted samples and handle missing observables fairly, (ii) three joint hooks (spectral, timing, and polarimetric) with explicit per-hook adjudication criteria, and (iii) an explicit screened-domain existence criterion.

We emphasize that the hook thresholds are treated as a specification rather than a post hoc fit target.

Accordingly, we adopt a single source of truth for the Tier-1 per-hook adjudication criteria (Table 1) and provide a lightweight execution recipe (Observer Quickstart) intended for archival screening.

Our aim is discriminability, not population inference. We do not claim universality or occurrence rates in the parent population.

This paper is not intended as a finalized microphysical theory or a population-wide prevalence analysis; it provides an auditable screened-domain protocol for determining whether the joint-hook pattern exists in a verifiable form.

We provide an audit-ready candidate-identification and existence-style screening protocol within a screened domain; questions of universality and prevalence are left to follow-up work.

3.1. *Observational hooks as an executable protocol*

Screens (apply first; to keep Tier-1 screening consistent).
Because the Tier-1 screening criterion is intentionally conservative, we apply minimal screens to define "vetted" high-spin, high-λ systems and to handle missing observables consistently:

A **Spin screen:** prioritize sources with an explicit spin estimate (e.g., Fe Kα reverberation/relativistic reflection, continuum-fitting where applicable, or polarimetric constraints).
For Tier-1 we treat $a_* \gtrsim 0.95$ as "high spin" and report the estimator and its systematic caveats.
We adopt $a_* \geq 0.95$ as an operational high-spin criterion, allowing margin for systematic uncertainties in spin estimators relative to the phenomenology-implied provisional activation scale ($a_* \simeq 0.97$), to be calibrated observationally.

B **Accretion screen:** define $\lambda \equiv L_{\rm bol}/L_{\rm Edd}$ using a documented $L_{\rm bol}$ method (SED-based where available) and a stated $M_{\rm BH}$ estimator; propagate conservative systematics (bolometric correction and virial-mass uncertainties).

C **Geometry/obscuration screen:** exclude Compton-thick or heavily absorbed cases where spectral curvature, lags, or polarization become instrumentally ambiguous.
If an observable is effectively inaccessible at useful S/N—i.e., it cannot yield a decidable per-hook adjudication under the fixed Tier-1 floors and criterion rule in Table 1 (given the stated uncertainties/limits)—treat it as *missing* rather than *failed*.

D **Compactness screen:** gate out pair-opaque states ($\ell \gtrsim 30$) where the hard component is suppressed.
Here the compactness screen acts as an observability gate: sources that violate the transparency/compactness requirement are treated as not testable for the Tier-1 hooks (screened), rather than as failed outcomes in per-hook adjudication.
Operationally, **screened** $\neq$ **failed**: Screen D withholds non-identifiable cases from Tier-1 bookkeeping rather than counting them as negative outcomes.
Accordingly, the Tier-1 existence-style claim is assessed only within the domain where the hooks are measurable; outside this domain the protocol makes no claim.

Screen D (compactness) is an identifiability gate, not a physical exclusion. Here "screened" means *information-limited* (observationally saturated), so such cases are withheld from Tier-1 bookkeeping.

This compactness cut should be read as an observational transparency gate, not as a claim that low-ℓ states are the default visible manifestation of SSPE.

If substantial dissipation already occurs close to the inner-flow/ergoregion interface, radiative transfer and pair opacity can readily push the system out of this transparent window.

Accordingly, the low-ℓ subset is best interpreted as a comparatively favorable observational window rather than as the generic state of the mechanism.

When the relevant Hook diagnostic is not decidable given the achieved bandpass/S/N (e.g., due to radiative-transfer or pair-opacity complications), we label it as *missing* rather than *failed*.

To prevent post-hoc reclassification, we predefine an operational floor for each Hook in Table 1; below this floor the outcome is *missing* by construction and does not enter Tier-1 adjudication.

This bookkeeping rule is fixed in advance and is applied identically to all objects in the vetted sample.

At high compactness, radiative transfer and pair opacity can saturate the 20–120 keV band and wash out timing/polarimetric diagnostics, so the Hook suite becomes *non-identifiable* rather than negative.

We therefore label $\ell > \ell_{\rm crit}$ sources as *screened*, under the present Tier-1 protocol, not as failures in Tier-1 bookkeeping.

In other words, Screen D prevents false negatives due to observational saturation, and it can be relaxed when higher-energy coverage or improved modeling becomes available.

This choice is conservative: it reduces the claim to the domain where the proposed screened-domain test is well-posed.

Tier definitions. We use Tier-1 to denote lightweight archival screening (existing hard-X spectroscopy and timing products with minimal assumptions).

We use Tier-2 for higher-cost follow-up and systematic tightening (e.g., dedicated polarimetry, broader-band coverage, and targeted re-analysis for sources near screening thresholds).

Hooks (apply jointly under Screens A–D). Operational definitions and per-hook adjudication criteria are given in Table 1; here we state only the physical intent:

Hook 1: Hard X/MeV shoulder. A hard excess/shoulder beyond a featureless baseline continuum.

The featureless baseline continuum is defined operationally as a shoulder-free continuum model fitted with the same background and line-of-sight absorption treatment as the shoulder-inclusive comparison model.

Hook 2: Lag hardening. A lag-energy relation consistent with harder photons lagging softer photons within an epoch (*epoch* as operationally defined in Appendix A).

Hook 3: Equator-aligned X-ray polarization. EVPA preferentially aligned with the disk plane, with a trend toward higher polarization degree at higher spin/high λ.

Robustness to the information-criterion threshold. We adopt a conventional threshold (fiducially ΔBIC (ΔAICc optional) ≥ 6) to keep the Tier-1 archival test operational and reproducible, but the protocol does not hinge on this single value.

In practice, we recommend reporting Hook-1 adjudication stability under a small bracket around the threshold (e.g., ΔBIC (ΔAICc optional) $\in [4, 10]$) as a sensitivity check; sources whose Hook-1 adjudication flips within this bracket are labeled ambiguous and deferred to Tier-2.

Lag convention. We adopt $E_{\rm ref} = 2$–4 keV in the rest frame unless stated otherwise; in the observed frame use $(E_{\rm ref}/(1+z))$ where feasible, and define $\tau > 0$ when the band at E lags the reference band.

In this convention, "lag hardening" corresponds to *hard lags*: $\tau(E)$ increases with energy.

(See Fig. 8 for a visual mapping of Screens and Hooks to the Tier-1/Tier-2 decisions.)

Observer Quickstart (Tier-1 archival screening; non-committal).

Step 1 (screen first):
Apply Screens A–D and document any missing observables.

Step 2 (Hook 1, spectral):
Perform the model comparison specified in Table 1 and report the preference metric.

Step 3 (Hook 2, timing):
Compute the lag diagnostic and report the criterion specified in Table 1.

Step 4 (Hook 3, polarization):
Test the EVPA alignment criterion specified in Table 1.

Step 5 (Tier-1 archival test):
If the vetted archival set recovers the joint-hook pattern in more than one source, the Tier-1 existence-style criterion is met and Tier-2 follow-up is motivated.
If zero or one such source remains, the claim is not established.
If screened or decidable coverage is too limited, the archival result is treated as inconclusive and carried forward to Tier-2 where appropriate.

Tier-1 archival test (screened-domain existence criterion).—Under Screens A–D, Tier-1 asks whether the SSPE-like joint-hook pattern is recovered in more than one source within the vetted high-spin/high-λ archival set.

If so, the Tier-1 existence-style criterion is met.

If zero or one such source remains, the claim is not established.

If screened or decidable coverage is too limited, the archival result is treated as inconclusive rather than positive.

Degeneracy control via joint coherence.—Each hook may admit degeneracies in isolation (e.g., complex coronal geometry or slim-disk radiative transfer can mimic shoulder-like curvature, and propagation/reverberation effects can shape energy-dependent lags).

SSPE therefore does not rely on any single signature in isolation; after applying Screens A–D, Tier-1 evaluates whether Hooks 1–3 jointly recover the expected pattern within the same source/epoch, with within-object model comparison (e.g., ΔBIC (ΔAICc optional) ≥ 6) applied where a competing spectral/timing description exists.

Minimal discrimination argument.—While each hook admits plausible degeneracies in isolation, reproducing the same-epoch co-occurrence of
(i) a statistically preferred hard-X/MeV shoulder,
(ii) a monotone hard-lag spectrum with $d\tau/d\log E > 0$, and
(iii) a hard-component EVPA preferentially aligned with the disk plane within the vetted high-a_*, high-λ subset typically requires fine-tuned coronal geometry and scattering configurations.

SSPE therefore does not claim uniqueness of any single signature; instead, it targets low-cost discriminability through joint observational coherence.

This design intentionally shifts the burden from microphysical closure to low-cost, operationally testable multi-observable consistency.

Single source of truth. We treat Table 1 as the authoritative Tier-1 specification.

Table 1. Operational hook specification

Item	Observable	Band / Product	Pass criterion	Primary confounders
Hook 1	Hard X/MeV shoulder	20–120 keV	ΔBIC (ΔAICc *optional*) ≥ 6	model degeneracy
Hook 2	Lag hardening	lag-energy relation	with $\tau > 0$ defined as a lag; pass if $d\tau/d\log E > 0$ and monotonic[b] within an epoch[a].	dilution
Hook 3	Equator-aligned X-ray polarization	2–10 keV hard component	hard-component EVPA $\parallel$ disk (equatorial plane)	jet contamination

Note— After applying Screens A–D, individual hooks may still retain residual degeneracies; SSPE is therefore evaluated by joint coherence across Hooks 1–3.
Screen D acts as an observability gate **(screened $\neq$ failed)**: sources outside the transparency domain are excluded rather than counted as failures.
Table 1 defines per-hook adjudication only, while Appendix A defines the object-level bookkeeping and the screened-domain existence criterion.

[a] An *epoch* denotes a contiguous analysis interval within which the source state is treated as approximately stationary for lag-energy estimation; the exact operational definition is given in Appendix A.

[b] Here "monotonic" means $\tau(E)$ is non-decreasing across the adopted energy bins within the quoted uncertainties (i.e., no statistically significant reversals).

4. POSITION IN CONTEXT

Summary.— SSPE is not a replacement for existing spin-extraction channels, but an equatorial complement. Thorne-limit arguments remain valid in their original thin-disk regime; BZ jets remain the natural outlet for MAD, strongly magnetized, radio-loud systems. SSPE instead targets the high-λ, high-spin, radio-quiet sector where neither tool fully constrains the phenomenology.

Classical spin-evolution work shows that radiation capture in a razor-thin Novikov–Thorne disk drives an asymptotic limit $a_* \simeq 0.998$ (K. S. Thorne 1974).

We regard this as a clean limiting case, not a universal ceiling: real quasars host geometrically thick, magnetized, outflow-bearing inner flows that violate the assumptions of the original derivation.

Likewise, the Blandford–Znajek (BZ) mechanism provides a robust path for launching powerful polar jets in magnetically arrested disks (e.g., R. D. Blandford & R. L. Znajek 1977; A. Tchekhovskoy et al. 2011), and can efficiently bleed spin energy in radio-loud, MAD-like systems.

However, BZ by itself does not explain the combination of high inferred spins, high Eddington ratios, and weak radio emission seen in many luminous quasars.

In this landscape, SSPE is best viewed as an *equatorial exhaust channel*.

When the ISCO approaches the ergoregion and the inner flow becomes thick and strongly coupled, an equatorially concentrated outflow component can transfer rotational energy back into the inner disk and cocoon.

This plausibly supports mildly to strongly super-Eddington luminosities while enforcing an effective "spin ceiling" $a_* \lesssim 0.97$–0.99 for systems that lack a powerful jet.

At the coarse-grained bookkeeping level adopted here, BZ and SSPE can be viewed as a see-saw: polar extraction tends to dominate MAD, radio-loud states, while equatorial extraction can become more prominent in high-λ, radio-quiet states.

The observational hooks and the screened-domain Tier-1 criterion in Sec. 3.1 are designed precisely to assess this equatorial channel.

In what follows, we use the shorthand "see-saw" to denote an anti-correlated power partition between the polar and equatorial channels, i.e., $P_{\rm BZ}$ increasing as $P_{\rm eq}$ decreases, at fixed $P_{\rm ext}$ (and fixed/small $f_{\rm oth}$).

5. DISCUSSION

5.1. *Synthesis across probes*

Taken in isolation, each observational tension motivating SSPE admits a plausible local explanation.

The core claim here is *not* that such explanations are impossible, but that their repeated co-occurrence in the same high-spin, high-λ, often radio-quiet subset makes a single shared inner-disk heating channel a lower-complexity hypothesis.

SSPE is offered as such a channel in a form that can be operationally assessed within the screened domain: after applying Screens A–D, the question is whether the joint-hook pattern defined in Sec. 3.1 emerges reproducibly in more than one source, rather than as a one-off anomaly.

These bounds align with Fig. 6–7 and are sufficient to reach $L/L_{\rm Edd} \sim 2$–3 with rarer 5–10× excursions.

Slim disks allow modest super-Eddington flows but no spin ceiling. BZ/MAD explain jet power yet not the radiative dominance of radio-quiet quasars.

This framework ties hyperluminous output and the spin ceiling via near-extremal, equatorial coupling.

Consistent with GRFFE/GRMHD studies, an equatorial return-current sheet can form within or near the ergoregion and become plasmoid-unstable; detached plasmoids can split into branches with $E - \Omega_{\rm H} L \lesseqgtr 0$ (captured vs. escaping), enabling rotational-energy extraction via a Penrose-like radiative channel (R. Penrose 1969; R. Penrose 2002; S. S. Komissarov 2004; W. E. East & H. Yang 2018; Z. Pan 2018; K. Parfrey et al. 2019; A. Bransgrove et al. 2021).

Relation to Blandford–Znajek jets. BZ is polar/Poynting-dominated; our mechanism is equatorial/radiative.

Hybrid states and an anti-correlated partition are expected; counterexamples (simultaneously strong jets and high radiative output) can occur in transitional MAD-like regimes.

Importantly, SSPE is not posed as a competitor to BZ, but as an equatorial channel that can coexist with the global magnetospheric topology required for BZ.

In this sense, SSPE may utilize the equatorial return current and inner-disk coupling present in BZ-favorable configurations in principle.

The relevant emission / deposition radius $R_{\rm eff}$ for the SSPE-driven equatorial outflow component is not known a priori: it depends on microphysics that current GRMHD and kinetic simulations have not yet fully resolved.

In this paper we therefore treat $R_{\rm eff}$ as a sensitivity parameter rather than a fitted quantity.

5.2. *Parameter space for equatorial-channel transparency*

Figure 6–7 illustrates how the Screen D $\ell = \ell_{\rm crit}$ carves the $(\epsilon_{\rm coup}, d)$ plane for two representative choices of $R_{\rm eff}$; the figure is meant as a diagnostic guide, not as a prediction for a specific object.

Interpreted conservatively, increasing $R_{\rm eff}$ enlarges the *allowed* transparent window, not the likelihood of direct escape.

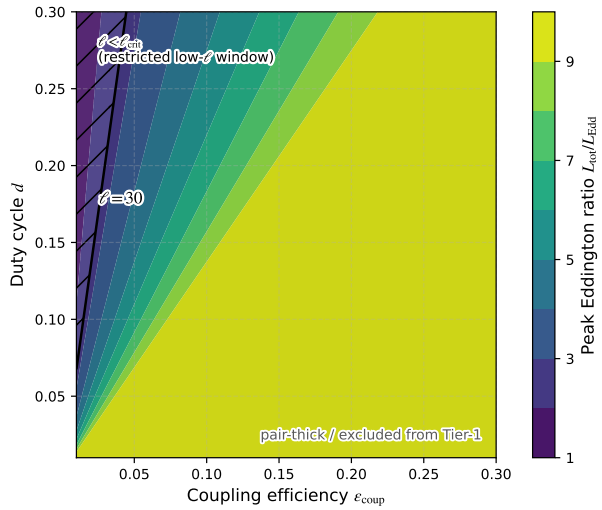

Figure 6. Restricted observational transparency window for the screened-domain test. Peak Eddington ratio $L_{\rm tot}/L_{\rm Edd}$ (color) on the $(\epsilon_{\rm coup}, d)$ plane; the black contour marks ℓ=30 from Eq. (6) with components defined in Eqs. (4,5). For $R_{\rm eff} = 100\, r_g$, the allowed region lies in the hatched low-ℓ area to the left of the $\ell = 30$ contour, where $L_{\rm tot}/L_{\rm Edd} \sim 2$–3 is typical with rarer 5–10× excursions; see App. A.

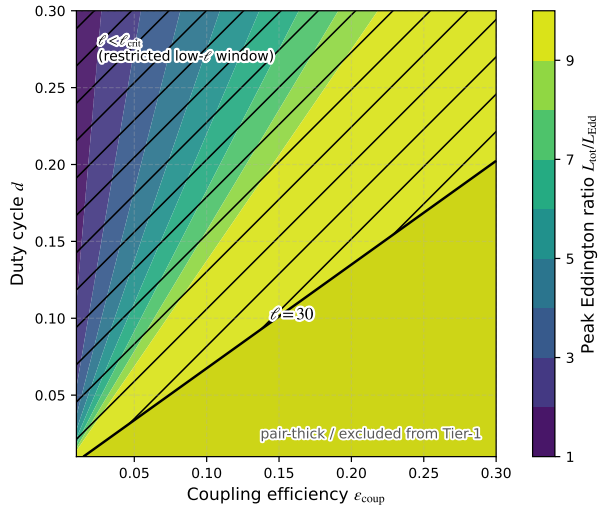

Figure 7. A larger dissipation footprint widens the restricted observational transparency window. Peak Eddington ratio $L_{\rm tot}/L_{\rm Edd}$ (color) on the $(\epsilon_{\rm coup}, d)$ plane for a larger footprint $R_{\rm eff} = 1000\, r_g$. The allowed region again lies in the hatched low-ℓ area to the left of the $\ell = 30$ contour (black; from Eq. (6) with components in Eqs. (4,5)), enlarging the viable window compared to Fig. 6.

5.3. *Power partition and spin budget*

Notation. We denote by $P_{\rm ext}$ the total power extracted from the Kerr spin reservoir (Sec. 2); it partitions into a polar Blandford–Znajek (BZ) channel, an equatorial channel, and residual channels:

$$P_{\rm ext} = P_{\rm BZ} + P_{\rm eq} + P_{\rm oth}, \tag{12}$$

$$P_{\rm BZ} = f_{\rm BZ} P_{\rm ext}, \quad P_{\rm eq} = f_{\rm eq} P_{\rm ext},$$

$$P_{\rm oth} = f_{\rm oth} P_{\rm ext}, \tag{13}$$

$$f_{\rm BZ} + f_{\rm eq} + f_{\rm oth} = 1. \tag{14}$$

Here $P_{\rm oth}$ collects channels not modeled in detail (e.g., pair-thick radiative outflows, advection, or dissipation that does not contribute to the equatorial hard component).

In the present work we focus on vetted, radio-quiet high-spin/high-λ systems where the hook-based program is designed to be most discriminating; a large $f_{\rm oth}$ is expected to weaken the spectral–timing–polarimetric hooks and can be treated within the screening and Tier-1 adjudication logic.

If independent evidence suggests that $f_{\rm oth}$ is not small or is rapidly variable, the see-saw diagnostic is treated as non-decidable for that object/epoch and is not used in the primary Tier-1 decision.

Importantly, $f_{\rm oth}$ does not modify the Tier-1 decision logic: for vetted objects with decidable joint-hook outcomes, failure to recover the required joint-hook pattern prevents that object from supporting the Tier-1 existence case, independent of how the extracted power is partitioned into residual channels ($P_{\rm oth}$).

A practical implication is that SSPE can affect systems that appear, by their large-scale phenomenology, to be BZ-favorable.

Even when the polar channel dominates the long-term power budget (i.e., $P_{\rm BZ} \gg P_{\rm eq}$ on average), a nonzero equatorial channel provides an internal heating pathway that acts directly on the disk photosphere rather than through external irradiation.

In such states, the observable tension is not "a jet plus a thin disk", but a jet coexisting with a disk that looks *too extended and too hot* for a strictly passive, radiatively efficient thin-disk interpretation—as suggested by microlensing/reverberation size inferences and wavelength-dependent lag amplitudes that exceed simple reprocessing expectations.

Within SSPE, the same bookkeeping that permits strong jet power also permits intermittent equatorial reheating episodes whose signatures should correlate with the spectral-timing hooks (Sec. 3) while remaining compatible with radio-loudness.

This is not an attempt to fit any single source; rather, it motivates a concrete null test: in jet-dominated quasars exhibiting anomalously large continuum sizes or overheated outer disks, the absence of *both* hard-component hooks (Hook 1/2) *and* disk-parallel hard-X polarization (Hook 3), after controlling for absorption and compactness screening, would disfavor equatorial feedback as the common cause.

This picture motivates looking for an anti-correlation between radio-jet dominance and equatorial reheating diagnostics within the high-a_*, high-λ subset, with hybrid states possible in transient MAD phases.

Our claim is modest: in the high-spin, high-Eddington subset, a non-zero equatorial contribution may at times be important or even dominant; in others, $f_{\rm BZ}$ may prevail. Outside this subset we make no claim.

Thin-disk tensions and coexisting channels.—Independent of the SSPE hypothesis, multiple lines of evidence have highlighted persistent tensions between thin-disk expectations and observations, including continuum-emitting sizes inferred from quasar microlensing and wavelength-dependent continuum lags that exceed simple reprocessing predictions (e.g., C. W. Morgan et al. 2010; Y.-P. Li et al. 2019; see also the review by E. M. Cackett et al. 2021 and the corona–disk coupling model of M. Sun et al. 2020).

Here $P_{\rm ext}$ denotes the net extractable power available outside the horizon; Eqs. (12)–(14) describe its internal bookkeeping, while the relation $P_{\rm ext} = P_{\rm BZ} + P_{\rm eq}$ is a coarse-grained *channel* partition (polar vs. equatorial) used only for the see-saw discussion.

Rather than forcing a single channel, our framework allows a *coexisting* phenomenological partition of magnetospheric power between a polar BZ-like outflow and an equatorial channel,

with $P_{\rm ext} \sim P_{\rm BZ} + P_{\rm eq}$ (or $f_{\rm BZ} + f_{\rm eq} \sim 1$) as a coarse-grained bookkeeping limit when $P_{\rm oth}$ is small.

For the see-saw discussion we often consider an effective two-channel view (polar vs. equatorial) at approximately fixed/small $f_{\rm oth}$, while the full bookkeeping remains Eqs. (12)–(14).

This motivates testing whether jet proxies (e.g., radio loudness) covary with disk-size/lag excess and with the equatorial-hook suite (Hooks 1–3), providing an observational handle on the proposed see-saw.

Testable prediction (radio–X joint anti-correlation). If the outgoing power partition approximately follows a see-saw relation ($f_{\rm BZ} + f_{\rm eq} \simeq 1$ at fixed $\dot{M}$ and spin), one may look for an anti-correlation between a jet proxy and an equatorial SSPE proxy within samples screened by Screens A–D.

Operationally, we propose comparing radio loudness (or core radio luminosity at fixed mass and λ) against the shoulder metric (ΔBIC (ΔAICc *optional*)) and/or the Hook 2 timing slope within matched bins of (M, a_*, λ).

A minimal operational statement concerns the sign: after matching, the rank correlation is expected to be negative ($\rho_S < 0$) with significance $p < 0.01$, if the proposed see-saw trend is present.

5.4. *Degeneracy with disk–corona phenomenology and how the protocol breaks it*

Disk–corona phenomenology can mimic individual hooks; our goal is to break degeneracy at the protocol level rather than by adjudicating microphysics.

(i) Spin ceiling as a nontrivial constraint. Standard phenomenological corona heating does not, by itself, enforce an upper envelope in the spin distribution.

SSPE ties the heating channel to black-hole rotational energy extraction and predicts an equilibrium spin ceiling via a torque/energy balance, providing an additional empirical constraint.

(ii) Same-epoch co-occurrence across distinct modalities. Even if a corona geometry can be tuned to mimic one hook, requiring the *same-epoch* (operational definition; Appendix A) co-occurrence of (1) a 20–120 keV shoulder with strong model preference, (2) lag hardening with monotone $d\tau/d\log E > 0$, and (3) disk-parallel polarization constitutes a low-dimensional joint constraint.

Reproducing all three simultaneously would typically require coordinated adjustments of multiple independent corona degrees of freedom across sources and epochs.

(iii) Black-hole linked scaling as a Tier-2 discriminator. SSPE predicts that the joint-hook strength (i.e., the significance of the joint-hook evidence as quantified by the available hook metrics for a given object/epoch; e.g., ΔBIC for Hook 1, $d\tau/d\log E > 0$ for Hook 2, and polarization degree for Hook 3 when available) should correlate with spin-linked quantities (e.g., the equilibrium spin trends and iron-line diagnostics), whereas a purely geometric corona interpretation lacks a generic mechanism to lock the same correlations without additional tuning.

In SSPE the heating channel is ultimately powered by black-hole rotational energy extraction; accordingly, we expect the same correlations at fixed $(M_{\rm BH}, \lambda)$.

A purely geometric multi-layer corona can reproduce individual observables, but it lacks a generic mechanism to lock the same correlations without additional tuning.

We therefore treat spin-linked scaling (e.g., via reflection/Fe K diagnostics; Fig. 8) as a Tier-2 discriminator; see the Tier-1/Tier-2 decision workflow in Sec. 3.1 and the Tier-2 timing extensions summarized in Table 2.

Section 5.5 discusses a speculative direct-leakage scenario; it is not required for the core protocol or Tier-1 adjudication logic.

5.5. *Rare direct quasi-beam events as ergoregion probes*

We stress that the following subsection is explicitly speculative and *not* required for the core SSPE framework: the primary observational program remains the hook-based spectral–timing–polarimetric tests in Sec. 3.1.

Normal outcome: reprocessing into Hooks 1/2. In the working picture adopted here, SSPE acts primarily as an *equatorial power channel*:

a narrow quasi-beam extracts rotational energy near the ISCO–ergoregion interface and re-injects it into the inner disk and surrounding cocoon.

Most of the time this energy is promptly reprocessed into the hard X/MeV shoulder (Hook 1) and a positive lag-energy slope ($d\tau/d\log E > 0$; Hook 2), with only modest direct leakage beyond the disk and torus.

These reprocessed components constitute the primary hooks discussed in Sec. 3.1.

This speculative leakage scenario is not required for the protocol and does not affect the Tier-1 decision bookkeeping summarized in Table 2.

Rare outcome: partial beam leakage. However, there is no guarantee that the inner flow is always geometrically and optically thick in the relevant directions.

In special configurations—for example, low H/R at the inner rim, local warps or gaps, or mild misalignment between the spin axis and the large-scale disk—

a fraction of the equatorial quasi-beam may find a comparatively transparent channel and emerge without full reprocessing.

In that case SSPE could produce *direct* signatures on timescales of a few r_g/c, much shorter than the thermal and viscous response of the disk.

Qualitative signature set. We emphasize that we do *not* attempt a full classification here, but the following features are natural in the present framework:

(1) **Ultra-fast, ultra-hard flares:** isolated spikes or bursts with variability timescales $\sim$few r_g/c, spectral peaks extending into the hard X-ray / the sub-MeV band (e.g., 0.2–1 MeV), and hardness ratios exceeding those of the quasi-steady shoulder.
(2) **Hard-to-soft evolution:** individual events in which the hardest photons lead a softer tail, consistent with a narrow beam depositing energy into a thicker reprocessing layer.
(3) **Transient polarization swings:** in principle, a direct beam component should exhibit strong polarization with EVPA tied to the local equatorial geometry; rapid changes in the effective beam direction could induce short-lived swings in hard X-ray polarization angle.

Candidate hunting grounds. At the level of working hypotheses, two regimes appear especially promising:

(i) **Lamppost fits with extremely compact source heights ($h \sim$ a few r_g).** In some bright quasars and Seyferts, lamppost spectral fits drive the inferred source height to values $h \sim$few r_g or lower, where the physical interpretation becomes ambiguous. Within SSPE such cases could instead correspond to equatorial quasi-beams emerging just above the horizon, rather than a compact source suspended on the spin axis.
(ii) **High-λ narrow-line Seyfert 1s (NLS1s).** These systems host relatively low-mass SMBHs accreting at high Eddington ratios and often display strong, rapid X-ray variability. A subset with very hard, short flares and strong high-energy excess would be natural laboratories for rare beam-leakage episodes.

We stress that these are *suggestive* classes, not firm identifications: detailed modeling and joint timing–spectral–

polarimetric analyses are required before any specific source can be claimed as a direct quasi-beam detection candidate.

Relationship to the main hooks. From the perspective of this paper, rare direct quasi-beam events are not an additional mandatory hook but a high-risk, high-reward extension.

The primary observational claim of the SSPE framework remains the *co-occurrence* of Hooks 1–3 under the screening criteria of Sec. 3.1.

A confirmed direct quasi-beam event would instead provide an exceptionally clean window into the ergoregion dynamics, placing much sharper constraints on the coupling microphysics than the present phenomenological treatment attempts to specify.

Scope (low-spin controls).— Signatures detected with comparable frequency in low-spin control samples (See Sec. 2.4) would indicate a spin-insensitive origin (general disk/corona physics) and are therefore *not* interpreted as SSPE in this work, without altering the protocol's Tier-1 decision bookkeeping.

5.6. *Validation Roadmap*

Observational fronts.

(i) Polarization vs. spin. Test that EVPA is equator-aligned and that the degree rises with a_*/λ; control for jet and scattering-cone geometries; report rank-correlation p-values and effect sizes.

(ii) Lag hardening. Within single epochs, verify that energy-resolved X-ray lags satisfy $d\tau/d\log E > 0$ and remain monotone across the usable band (relative to a stated $E_{\rm ref}$); mask absorption events; provide cadence/noise tests. Tier-2: where available, check consistency with X→UV/optical reprocessing lags.

(iii) Spectral compactness screen. Use cutoff-energy/pair proxies to flag high-ℓ cases (non-transparent) and separate them from the transparent sample used for hooks.

(iv) BZ–equatorial power partition (see-saw). Using Sec. 5.3 as the bookkeeping ($P_{\rm ext} = P_{\rm BZ} + P_{\rm eq} + P_{\rm oth}$; $f_{\rm BZ} + f_{\rm eq} + f_{\rm oth} = 1$), test whether states (or sources) with larger jet-dominance proxies (e.g., radio loudness, core dominance, or independent jet-power estimates) systematically show weaker equatorial reheating diagnostics (Hooks 1–3), at fixed spin/λ and after screens.

Report partial correlations (controlling for inclination and absorption) and state-by-state transitions where available.

This see-saw test is a cross-check on the power partition, not an additional hook in the primary Tier-1 observational program.

Simulation fronts.

(i) GRMHD near the ergoregion. Resolve the equatorial return-current sheet; quantify energy partition into positive/negative branches ($E - \Omega_{\rm H} L \gtrless 0$); compare with torque-balance trends.

(ii) Radiative transfer with pairs. Propagate hard X/MeV emission through compactness-limited media; map when the shoulder survives vs. is quenched (link to ℓ cut).

(iii) Population modeling. As a follow-up exercise beyond the present Tier-1 scope, combine spin distributions with transparency and coupling priors to estimate the fraction of sources expected to satisfy the joint-hook requirement.

Table 2. Validation roadmap matrix. What can be tested now, and how each hook informs the screened–domain Tier-1/Tier-2 program.

Hook	Minimal data products	Executable now?	Interpretation if absent after screens
Hook 1	Hard X-ray spectrum (background-controlled) + model comparison (ΔBIC (ΔAICc *optional*) ≥ 6)	**Yes**	If absent after screens: featureless continua are favored over a hard shoulder in the vetted bright subset, weakening the case for Hook 1 in Tier-1/Tier-2 follow-up
Hook 2	Energy-resolved X-ray light curves (multiple bands; sufficient counts) + lag-energy spectrum $\tau(E)$ (relative to a stated $E_{\rm ref}$) Tier-2: X→UV/optical reprocessing lags where available	**Partly**	If absent after screens: lag-energy slopes fail to show positive, monotonic hardening within an epoch[a] in the vetted bright subset, weakening the case for Hook 2 in Tier-1/Tier-2 follow-up
Hook 3	Hard-component X-ray polarimetry + EVPA reference to disk plane + spin/λ proxies	**Limited**	If absent after screens: EVPA shows no preference for equatorial alignment and no trend with inferred spin/λ, weakening the case for Hook 3 in Tier-1/Tier-2 follow-up

NOTE—The primary Tier-1 criterion used in this work is the joint-hook rule within the screened domain. Table 2 is a validation roadmap, not the single source of truth for Tier-1 bookkeeping.

[a] Hook 2 is marked as *Partly* mainly because simultaneous, high-time-resolution multi-band timing datasets are still rare (e.g., strict epoch matching across bands).

Primary (Tier-1) criterion (screened-domain existence-style). Tier-1 does not perform population inference.

Instead, after applying Screens A–D, it asks whether the SSPE-like joint-hook pattern is recovered in more than one vetted source in the archival sample.

If two or more vetted sources satisfy the joint-hook pass condition, the Tier-1 existence-style condition is satisfied.

If *zero or one* such source remains, the existence-style claim is not established.

If the number of vetted or decidable cases is too small, or if missing outcomes dominate because of data limitations, the archival Tier-1 result is treated as inconclusive rather than negative.

We define an object as decidable for Tier-1 bookkeeping if and only if at least two hooks yield non-missing outcomes (pass/fail) within the chosen representative epoch.

Each object is counted once (using a single representative epoch per object to avoid double-counting), and missing outcomes are excluded by construction from decidable Tier-1 bookkeeping.

Operational criterion. The executable screened-domain existence criterion is defined in Appendix A and is the rule used for Tier-1 bookkeeping in this work.

Here Tier-1 denotes the executable archival-screening and bookkeeping layer used in this work, whereas Tier-2 denotes non-decisive follow-up considerations requiring future targeted observations or improved sensitivity.

Secondary (Tier-2) supporting patterns (non-decisive). The following patterns are not part of the executable Tier-1 criterion and are not used for Tier-1 bookkeeping.

They are included only as secondary contextual considerations for future Tier-2 tests and broader interpretation.

Beyond the primary Tier-1 criterion, the following patterns, while individually non-decisive and potentially confounded, would weaken the broader SSPE interpretation in that subset:

1. Systematic absence of the Hard-X/MeV shoulder (20–120 keV, obs), together with systematically weak high-ionization-line signatures relative to a matched control sample.
2. Reverberation lags from X→UV showing no energy dependence (no inward reheating signature).

If *two or more* items above hold simultaneously for the same vetted high-spin, high-λ subset, they would weaken the broader motivation for invoking an equatorial SSPE-like channel in that subset.

6. CONCLUSION

Self-sustained Penrose excitation near extremal Kerr SMBHs can inject rotational energy into the accretion disk, yielding sustained 2–3× Eddington with rarer 5–10× episodes while enforcing a spin ceiling.

Accordingly, the present contribution is best read as a pathfinding observational protocol: a falsifiable, screened-domain route for testing whether this specific joint-hook pattern exists, rather than as a completed microphysical closure.

APPENDIX

A. OPERATIONAL PROTOCOL (SELF-CONTAINED)

Screens A–D (apply first; see Sec. 3.1 for details; summarized here).
Operational floors and pass/fail thresholds are defined in Table 1; below the floor each hook is treated as missing by construction.
Screen A (spin): require an available spin proxy; adopt vetted high-spin as $a_* \gtrsim 0.95$.
Screen B (λ): require an Eddington-ratio estimate ($\lambda \equiv L_{\rm bol}/L_{\rm Edd}$) with stated systematics; select the high-λ subset.
Screen C (geometry/obscuration): exclude Compton-thick/heavily absorbed cases and extreme inclination; if a hook is plausibly blocked at useful S/N—i.e., the observable is **non-decidable** as pass/fail under the fixed Tier-1 floors and decision rule in Table 1 (given the stated uncertainties/limits)—treat it as *missing* rather than failed.
Screen D (compactness): require $\ell < \ell_{\rm crit}$ (fiducial $\ell_{\rm crit} = 30$); pair-opaque states are *screened* (not failures).
Compactness ℓ is defined as in Eq. (6).
In operational use, this screen is applied using compactness-related proxies independent of the Hook 1 pass/fail decision (e.g., cutoff-energy/pair-opacity diagnostics or other stated estimates of $L_{\rm self}$ and $R_{\rm eff}$).
If these proxies do not permit a defensible adjudication of transparency under the fixed floors, the object/hook is treated as non-decidable (missing) rather than failed.
Absence of a visible Hook 1 feature is not by itself sufficient to classify a source as low-ℓ transparent; possible quenching at high compactness must be excluded independently before Hook 1 is counted as failed.
Lag convention (used in Hook 2). We measure $\tau(E)$ relative to a fixed reference band $E_{\rm ref}$ (adopt $E_{\rm ref} = 2$–4 keV in the rest frame; use $(E_{\rm ref}/(1+z))$ in observed frame where feasible); define $\tau > 0$ when the band at E lags the reference band.

Hooks (apply jointly under Screens A–D).
Hook 1 (shoulder): passes if ΔBIC (ΔAICc *optional*) ≥ 6 in the 20–120 keV observed band.
Hook 2 (lag hardening): passes if $d\tau/d\log E > 0$ and $\tau(E)$ is monotone within an *epoch*.
If the achieved counts/S/N within an accepted epoch are insufficient to determine the sign of $d\tau/d\log E$ under the adopted binning and uncertainties, Hook 2 is treated as missing rather than failed.
Hook 3 (equatorial polarization): EVPA $\parallel$ disk/equator, with polarization degree increasing with spin/λ.

Epoch definition (operational). We define an *epoch* as a single observing block during which (i) the hardness ratio and (ii) the broadband flux, computed with fixed band definitions and a fixed reduction/selection procedure within the block, are statistically consistent with constancy.
Operationally, epoch constancy requires

$$\frac{|HR_i - HR_0|}{\sigma_{HR,i}} \;\leq\; 2 \quad \text{and} \quad \frac{|F_i - F_0|}{\sigma_{F,i}} \;\leq\; 2, \tag{A1}$$

over the block, where (HR_0, F_0) are the block means and $(\sigma_{{\rm HR},i}, \sigma_{F,i})$ are the 1σ uncertainties of the individual measurements (not the standard error of the block mean), including instrumental noise.

In practice, Eq. (A1) is evaluated pointwise for the in-block measurements under the fixed reduction/selection procedure, and the epoch is accepted if at least 95% of points satisfy the inequality for both HR and flux (ensuring robustness to occasional outliers).

If this criterion is not met, the hook is treated as *missing* for that epoch (not failed).

Counts and units (Tier-1). Tier-1 bookkeeping is performed at the object level within the screened domain.

N_{total}: number of objects in the vetted sample.
Vetted sample.—In Tier-1, N_{total} counts only objects that pass Screens A–D; objects screened out by any screen (including Screen D) are not included in N_{total}.
N_{miss}: number of objects for which the joint-hook rule is non-decidable because fewer than two hooks yield adjudicable outcomes under Table 1.
$N_{\mathrm{eff}} \equiv N_{\mathrm{total}} - N_{\mathrm{miss}}$: effective number of objects with decidable Tier-1 outcomes.
N_{pass}: number of objects in the effective sample that satisfy the joint-hook requirement.

Single source of truth (Tier-1). Tier-1 is an existence-style gate within the screened domain, not a population-inference test.

$$N_{\mathrm{pass}} \geq 2, \tag{A2}$$

which is the executable Tier-1 condition used in this work: the SSPE-like joint-hook pattern must be recovered in at least two vetted objects within the screened domain.

This threshold is not a significance test or a population-level estimator; it functions as a minimal anti-one-off gate for an existence-style archival protocol.

A single pass leaves only a suggestion-level result, whereas two or more independent passes satisfy the existence-style criterion within the screened domain.

If $N_{\mathrm{pass}} \leq 1$ after the pre-registered screens and bookkeeping, the Tier-1 existence-style condition is not met.

If the decidable sample is too small or too incomplete for Tier-1 adjudication, the Tier-1 result is treated as inconclusive rather than negative.

Note (data quality). If a Tier-1 hook cannot be adjudicated as pass/fail at useful S/N under the fixed floors in Table 1 (given the stated uncertainties/limits), its outcome is treated as missing by construction (not failed) and excluded from Tier-1 bookkeeping.

If the lag-energy slope is measurable under the fixed floors and is consistent with zero or negative, Hook 2 is counted as failed, not missing.

Hook 2 is treated as missing only when the data quality is insufficient to adjudicate the slope/monotonicity at the adopted floors.

Decision (Tier-1 bookkeeping). Tier-1 adjudication follows Eq. (A2) together with the missing-by-construction rule defined above.

Toy mapping (optional visual aid). See Fig. 8.

Toy mapping (optional visual aid).
The figure below provides an illustrative mapping from the equilibrium spin $a_{\rm eq}$ to two observables discussed in the hook protocol. It is a visual aid only and is not used in Tier-1 bookkeeping.

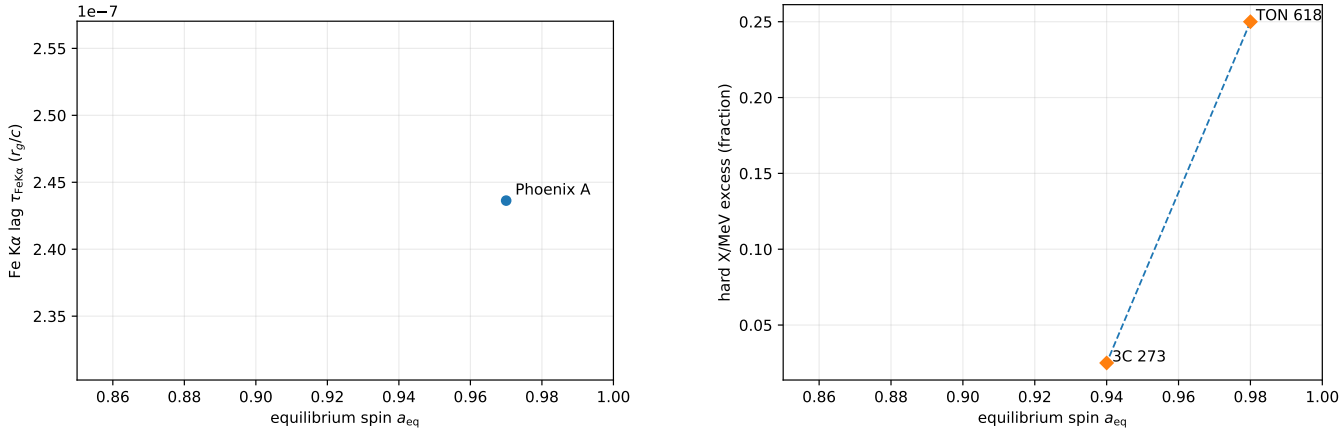

Figure 8. Hook observables vs. equilibrium spin.
Left: Fe Kα reverberation lag $\tau_{\rm FeK}$ (in r_g/c) as a function of $a_{\rm eq}$; lags increase toward lower $a_{\rm eq}$ (stronger coupling).
Right: Illustrative fraction $f_{\rm hook}$ of vetted objects expected to show a given hook within the toy mapping (not the Tier-1 bookkeeping fraction) versus $a_{\rm eq}$: the fraction rises steeply toward low $a_{\rm eq}$.
These figures summarize how the equilibrium spin $a_{\rm eq}$ maps onto two directly testable observables discussed in the hook protocol. See Sec. 3.1 and App. A for the operational definitions and pass criteria.

CAUTIONARY NOTE.

The claims advanced in this work are modular and falsifiable. Alternative mechanisms may dominate in other source classes.

AUTHOR CONTRIBUTIONS

Conceptualization, modeling, analysis, visualization, and writing: J. Wakabayashi.

COMPETING INTERESTS

The author declares no competing interests.

DATA AND CODE AVAILABILITY

All figures can be regenerated from scripts in the accompanying repository; data sources and acquisition steps are documented in the README.
An archived OSF snapshot (Version 1.1) is available: https://doi.org/10.17605/osf.io/q7dth.

COMMUNICATION AND MEDIA

Media note: This paper and its reproducibility package are the sole authoritative sources.

ACKNOWLEDGMENTS

This work stands on decades of insight into black-hole accretion, spin, and energy extraction. I am indebted to the community that built the modern framework of quasar physics—from classical Penrose energy extraction and disk theory to spin-jet coupling and polarimetry—and to teams who made public data and tools available.
Use of large language models. Large language model assistants (Google Gemini and OpenAI ChatGPT) were used for editorial support, including clarity editing, formatting suggestions, and figure-placement/LaTeX troubleshooting. No novel data, equations, or results were produced by these tools. No confidential or unpublished data were provided to them. All analysis, derivations, and conclusions are by the author, who takes full responsibility for the content; these tools are not authors.
I also thank colleagues and readers who provided critical comments on early drafts. Any remaining errors are mine.

Software: `latexmk`, `AASTeX701`

REFERENCES

Abramowicz, M. A., Czerny, B., Lasota, J.-P., & Szuszkiewicz, E. 1988, The Astrophysical Journal, 332, 646, doi: 10.1086/166683

Blandford, R. D., & Znajek, R. L. 1977, Monthly Notices of the Royal Astronomical Society, 179, 433, doi: 10.1093/mnras/179.3.433

Bransgrove, A., Ripperda, B., & Philippov, A. 2021, Phys. Rev. Lett., 127, 055101, doi: 10.1103/PhysRevLett.127.055101

Cackett, E. M., Bentz, M. C., & Kara, E. 2021, iScience, 24, 102557, doi: 10.1016/j.isci.2021.102557

East, W. E., & Yang, H. 2018, Physical Review D, 98, 023008, doi: 10.1103/PhysRevD.98.023008

Gralla, S. E., & Jacobson, T. 2014, Monthly Notices of the Royal Astronomical Society, 445, 2500, doi: 10.1093/mnras/stu1690

Gu, L., et al. 2025, A&A, 704, A146, doi: 10.1051/0004-6361/202557189

Komissarov, S. S. 2004, Monthly Notices of the Royal Astronomical Society, 350, 427, doi: 10.1111/j.1365-2966.2004.07672.x

Li, Y.-P., Yuan, F., & Dai, X. 2019, MNRAS, 483, 2275, doi: 10.1093/mnras/sty3245

Lightman, A. P., & Zdziarski, A. A. 1987, ApJ, 319, 643

Morgan, C. W., Kochanek, C. S., Morgan, N. D., & Falco, E. E. 2010, ApJ, 712, 1129, doi: 10.1088/0004-637X/712/2/1129

Pan, Z. 2018, Physical Review D, 98, 043023, doi: 10.1103/PhysRevD.98.043023

Parfrey, K., Philippov, A. A., & Cerutti, B. 2019, Physical Review Letters, 122, 035101, doi: 10.1103/PhysRevLett.122.035101

Penrose, R. 1969, Nuovo Cimento Rivista Serie, 1, 252

Penrose, R. 2002, General Relativity and Gravitation, 34, 1141, doi: 10.1023/A:1016578408204

Sun, M., Xue, Y., Brandt, W. N., et al. 2020, Astrophysical Journal, 902, 7, doi: 10.3847/1538-4357/ab789e

Svensson, R. 1984, MNRAS, 209, 175, doi: 10.1093/mnras/209.2.175

Tchekhovskoy, A., Narayan, R., & McKinney, J. C. 2011, MNRAS, 418, L79, doi: 10.1111/j.1745-3933.2011.01147.x

Thorne, K. S. 1974, Astrophysical Journal, 191, 507, doi: 10.1086/152991